



COVER: intended masks versus the delivered task

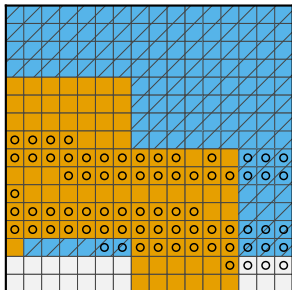
(a) Legacy: before collation



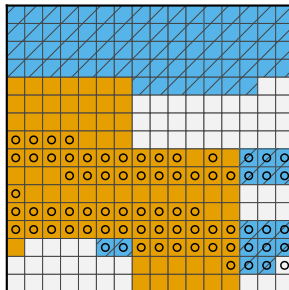
(b) Legacy: exact final indices



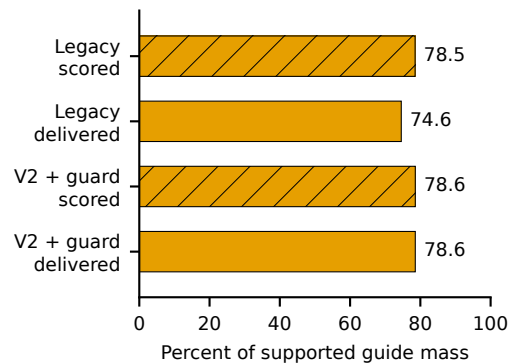
(c) V2: scored / sampled indices



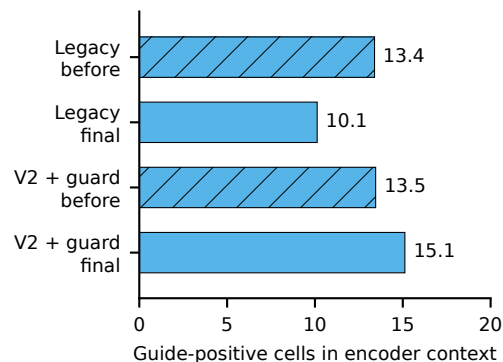
(d) V2 + guard: exact final indices



(e) Mean hidden guide mass
573 valid views, full-size batches



(f) Mean visible tissue
576 views, full-size batches



Target union
 Encoder context
 Guide-positive tissue
 Neither

Maps: one fixed Training view, 16 × 16 token grid; no raw OCT pixels. Charts: cohort means.

V2 scores delivered prefixes; its final context guard is a separate guide-aware intervention.