
Location Beats Shape: A Controlled Comparison of Anatomy-Guided Masking for Joint-Embedding Predictive Pretraining on Retinal OCT

Anonymous Author(s)

Affiliation

Address

email

Abstract

Masked predictive pretraining asks a model to reconstruct the representation of hidden image regions. Which regions are hidden is a free design choice, and in medical imaging there is an intuitive answer: hide the tissue that carries the diagnosis. We test that intuition under tightly controlled conditions. Six masking policies are continued from a single shared I-JEPA checkpoint on 3D retinal OCT, differing *only* in how predictor targets are placed, and evaluated with a frozen linear-probe protocol for glaucoma classification (FairVision, $N=3000$). Guidance does help: restricting targets to retinal tissue improves AUC by +0.0120 over unguided masking at the matched epoch (95% CI [+0.0068, +0.0173]), and a band located by a first-order intensity statistic improves it by +0.0109 at epoch 100 ([+0.0057, +0.0160], $p<0.0001$). All six contrasts against the null, at all three epochs, exclude zero. But anatomical *precision* does not help. A policy that places 97% of masked patches on tissue, using a trained retinal segmentation model to shape the targets, is *indistinguishable from unguided masking* at the matched epoch (+0.0013, CI [−0.0055, +0.0082]), as is a coverage-constrained policy (+0.0002) — while the two policies that merely *place* rectangles on tissue gain an order of magnitude more. Carried past epoch 50 the coverage-constrained policy does not merely stall, it falls *behind* unguided masking (−0.0084 at epoch 75, CI [−0.0146, −0.0026]), so over-constraining where the model may look is worse than not guiding it at all. Direct measurement of mask geometry across 600 slices shows the fraction of anatomy hidden is uncorrelated with downstream AUC. The best policy consults no segmentation model at all. A subgroup audit over 20 probes finds the same worst-served group every time; gains reach every group and every disease stage, mild disease most of all (+0.0137), but the difficulty ordering is untouched: the severe-to-mild gap stays within 0.0114 of constant across every policy. Because each policy was pretrained once, we report these as single-run observations rather than an established ranking.

1 Introduction

Self-supervised learning has become the default route to representation quality in medical imaging, where expert annotation is expensive and pathology is rare [Azizi et al., 2021, Zhou et al., 2023, Qiu et al., 2024]. Joint-embedding predictive architectures such as I-JEPA [Assran et al., 2023] learn by predicting the *representation* of masked target blocks from a visible context block, avoiding pixel reconstruction and its bias toward high-frequency detail [He et al., 2022, Xie et al., 2022].

I-JEPA samples its targets geometrically: rectangles placed uniformly at random. In natural images this is defensible, since informative content is spread broadly across the frame. In a retinal OCT B-scan it is not obviously so. A large fraction of the image is dark vitreous and background; the diagnostic signal for glaucoma is concentrated in a thin band of retinal tissue. This motivates an intuitive hypothesis, one that several recent methods adopt [Li et al., 2022, 2024, Wang et al., 2025, Shi et al., 2022]:

Direct the predictive objective at clinically meaningful anatomy rather than at arbitrary image regions.

We test this hypothesis directly and at several strengths. Rather than proposing one anatomy-aware method and reporting that it beats a baseline, we construct a *ladder* of masking policies of increasing anatomical specificity — from unguided rectangles, through rectangles constrained to lie on tissue, to ragged targets whose shape follows the segmented retina — and hold everything else fixed. All arms continue from one shared pretrained checkpoint, use identical optimiser, schedule, and effective batch size, and are evaluated under one frozen-probe protocol.

The result contradicts the intuition in an informative way. Guidance helps, but the benefit does not increase with anatomical precision — it disappears. The two policies whose targets track the segmented retina most faithfully are statistically indistinguishable from unguided masking at the matched epoch, while two policies that merely place ordinary rectangles on tissue gain an order of magnitude more. The best performer uses no segmentation model at all, locating a band by a per-column intensity centroid.

Contributions.

- A controlled six-arm study isolating masking policy in medical SSL, with a shared ancestor checkpoint and a single evaluation protocol (21 frozen probes).
- Evidence that guidance helps but anatomical precision does not, including the specific control — rectangles restricted to retinal tissue — that separates “avoid background” from “target anatomy”.
- Direct measurement of mask geometry for every policy on 600 slices, showing that fraction-of-anatomy-hidden does not order downstream AUC, and an explicit accounting of the axes on which the anatomy arms are confounded.
- A subgroup audit and a 10^{-5} reproduction of the headline result from public weights on different hardware and precision.

2 Related work

Masked image modelling and JEPAs. Masked autoencoders reconstruct pixels [He et al., 2022, Xie et al., 2022]; BEiT [Bao et al., 2022] and iBOT [Zhou et al., 2022] predict discrete tokens; data2vec [Baevski et al., 2022] and I-JEPA [Assran et al., 2023] predict latent representations. Video and 3D extensions follow the same template [Bardes et al., 2024, Chen et al., 2023b, Taleb et al., 2020]. Recent analyses question how much of the benefit derives from the masking distribution itself [Balestrieri and LeCun, 2025, He et al., 2025].

Informed masking. A substantial line of work argues that *what* is masked matters. ADIOS [Shi et al., 2022] learns an adversarial masking model; SemMAE [Li et al., 2022] uses learned semantic parts; AttMask [Kakogeorgiou et al., 2022] masks by attention; HardPatches [Wang et al., 2023a] and AutoMAE [Chen et al., 2023a] target difficulty. In the medical domain, AnatoMask [Li et al., 2024], AnatomyMAE [Ceballos Arroyo et al., 2026], MSMAE [Mao et al., 2023], and VasoMIM [Huang et al., 2026] explicitly steer masking toward anatomy, and Wang et al. [2025] argue directly for masking clinically relevant regions. Our study is a controlled test of that family’s central premise on one task, and finds it does not hold in the direction usually assumed.

Evidence that informed masking is not uniformly better. The picture is less one-sided than the paragraph above suggests. The original masked-modelling papers already report that random sampling is hard to beat: MAE’s ablation puts random-75 at 73.5 linear against block-75 at 63.9 [He

et al., 2022], and SimMIM finds random beats its best block variant [Xie et al., 2022]. Among informed schemes, AttMask sits within 0.2 points of random [Kakogeorgiou et al., 2022], AutoMAE ties MAE under full fine-tuning [Chen et al., 2023a], HPM gains only when randomness is retained and *loses* when restricted to the hardest patches [Wang et al., 2023b], and attention-guided MAE underperforms MAE in the low-label regime [Sick et al., 2025]. Most directly, SemMAE’s own part-quality ablation shows that *imperfect* semantic parts add nothing over no parts at all [Li et al., 2022], and Guo et al. [2025] reproduce SemMAE below random MAE under a common protocol, with the deficit largest under frozen evaluation — also our protocol. Chen et al. [2023a] name the mechanism a “patch selection dilemma”: a slight foreground bias helps, but raising it too far starves the encoder of the evidence it must predict from. Conversely, the closest positive precedent for our best policy uses no segmentation model either: Lee et al. [2025] select foreground from normalised intensity alone and gain across five 3D medical benchmarks. Appendix H tabulates these.

We therefore do *not* claim to overturn the informed-masking literature. That random masking is a strong baseline, and semantic guidance wins only sometimes, is known and contested. What we add is a controlled measurement of where a *trained medical segmentation model* falls on that spectrum when everything else is held fixed, in a setting — 3D retinal OCT, joint-embedding prediction, frozen probe — where we could find no prior controlled comparison.

Retinal foundation models and OCT. RETFound [Zhou et al., 2023], VisionFM [Qiu et al., 2024], URFound [Yu et al., 2024], and OCTCube [Liu et al., 2026] pretrain on large retinal corpora; OCT-specific masked modelling is explored by Pissas et al. [2024]. MIRAGE [Morano et al., 2025] provides multimodal OCT segmentation and supplies the anatomical guidance used by several of our arms. Deep glaucoma detection from OCT is established [Maetschke et al., 2019, Ran et al., 2019, Noury et al., 2022].

Fairness in medical imaging. Performance disparities across demographic groups are widely documented [Seyyed-Kalantari et al., 2021, Larrazabal et al., 2020, Gichoya et al., 2022]. FairVision [Luo et al., 2024a] and Harvard-GF [Luo et al., 2024b] provide OCT data with demographic attributes; FairMedFM [Jin et al., 2024] and Shi et al. [2025] benchmark fairness for medical foundation models. We report a subgroup analysis as a secondary finding.

3 Method: a ladder of masking policies

3.1 Background

I-JEPA partitions the token grid of an image into a context block and M target blocks. A context encoder f_θ embeds the visible tokens; a predictor g_ϕ maps that embedding plus positional queries to predicted representations of the target tokens; an exponential-moving-average target encoder $f_{\bar{\theta}}$ supplies the regression targets. With a 16×16 token grid the design choice we vary is which of the 256 cells become targets.

3.2 The six policies

All policies produce $M=4$ target blocks and one context block under identical scale and aspect-ratio ranges. They differ only in placement and shape.

RANDOM (null). Stock I-JEPA multiblock. Four rectangles placed uniformly at random anywhere in the grid. No image content is consulted.

ENVELOPE (random-within-retina). Identical rectangles — same shape, size and count as RANDOM — but rejection-sampled onto the retinal envelope predicted by MIRAGE [Morano et al., 2025]. Only *location* changes. This is the control that separates “stop wasting targets on background” from “target anatomy specifically”.

INTENSITY (segmentation-free). A band whose vertical position is located per slice by a per-column intensity-weighted row centroid, smoothed across columns. It adapts to the retina’s position and curvature using a first-order intensity statistic, and consults **no segmentation model**. It uses no ground-truth annotation of any kind. (Our released code and checkpoints label this policy `oracle`, a historical name from when it was intended as an upper bound.

Six masking policies on one B-scan, drawn with the production samplers.
Anatomical specificity increases left to right, top to bottom.




Figure 1: The six masking policies on one identical B-scan, drawn with the production samplers. Each predictor target block has its own colour; the green contour is the anatomy reference. Anatomical specificity increases along the ladder: ENVELOPE and COVER place rectangles on tissue, while the anatomy arms grow targets that trace the retina almost exactly. Downstream AUC does not increase with that specificity (Table 1). Figure 4 repeats this over five B-scans spanning the volume.

132 We rename it here because it is neither an oracle nor an upper bound: it reads only the input
133 image.)

134 **ANATOMY-V1 / ANATOMY-V2.** Ragged connected components grown on the MIRAGE score map,
135 so target *shape* follows tissue rather than being rectangular. v2 additionally bridges diagonal
136 adjacencies.

137 **COVER f .** Targets chosen greedily to cover the anatomy subject to a hard floor leaving a fraction f
138 of tissue visible in the context. We pretrain $f=0.21$.

139 Figure 1 shows all six on the same B-scans.

140 3.3 Hypotheses

141 The ladder yields three testable statements:

142 **H1** *Avoiding background helps.* ENVELOPE > RANDOM.

143 **H2** *Anatomical precision helps further.* anatomy arms > ENVELOPE.

144 **H3** *Any benefit is due to anatomical targeting, not to incidental changes in mask area, target count,*
145 *or context size.*

146 4 Experimental setup

147 **Data.** FairVision [Luo et al., 2024a] glaucoma OCT volumes. Each volume is a 3D scan resampled
148 to 100 B-scans of 256×256 after transform. The label is a binary glaucoma diagnosis. Splits are
149 fixed and patient-disjoint as released; downstream evaluation uses a held-out test split of $N=3000$
150 volumes (1466 positive / 1534 negative), seed 42, fixed loader order. No test volume is seen during
151 pretraining, probe fitting, or model selection: the probe head is selected on a separate validation
152 split. Labels are byte-identical across arms, so any two arms are paired-comparable on the same
153 cases. The dataset carries self-reported demographic attributes, which we use only for the subgroup
154 analysis in Section 5.3 and never as model inputs.

Pretraining. Patch-level I-JEPA, ViT-Base, 256×256 crops, patch size 16, predictor depth 6, EMA target encoder. **Every arm resumes from the same epoch-25 checkpoint** and continues with identical optimiser, learning-rate and weight-decay schedules, and effective batch size (512 for all arms). Guidance is ramped linearly from epoch 25 to 30 and held at full strength thereafter. The *only* difference between arms is the mask sampler.

Arms differ in how far they were carried, and we do not pretend otherwise. RANDOM, INTENSITY and ENVELOPE reach epoch 100. ANATOMY-V2 has valid probes to epoch 50 (later checkpoints follow a target-precision splice and are excluded, Appendix D). ANATOMY-V1 has only epoch 30. COVER $f=0.21$ was deliberately halted at epoch 73. Epoch 50 is therefore the only epoch at which five policies are simultaneously available, and it carries every cross-policy comparison that involves the anatomy or coverage arms.

Evaluation. The encoder is frozen. Per-slice features are mean-pooled over the volume, followed by LayerNorm and a linear head. The head is selected by validation AUC and reported on the held-out test split. We report paired bootstrap confidence intervals over test volumes (10,000 resamples, stratified by class, the same resampled index set applied to every arm) and DeLong tests for correlated ROC curves [DeLong et al., 1988], both of which are valid here because all arms are evaluated on the same cases in the same order. Each test volume is a distinct subject, so case-level resampling is already subject-level and no cluster correction is required [Obuchowski, 1997, Rutter, 2000]. We follow Maier-Hein et al. [2024] in reporting the operating-point metrics alongside AUC rather than AUC alone (Appendix A).

Numerical precision, and how we handle it. The probe harness defaults to mixed precision. The RANDOM, INTENSITY and ENVELOPE long-horizon probes were fitted under that default (fp16); the ANATOMY-V1, ANATOMY-V2, COVER and ancestor probes were fitted with autocast explicitly disabled (fp32). We therefore do *not* claim a single protocol across every probe in the study. Instead we partition the comparisons: all headline contrasts in Table 1 are drawn from arms sharing both epoch *and* precision, and any contrast that would cross the precision boundary is marked as such and excluded from headline claims. Section 5.4 reports a direct measurement of how large the precision effect actually is.

5 Results

What these numbers can and cannot establish. Each policy was pretrained exactly once, so policy is perfectly confounded with one stochastic optimisation path after the fork. The intervals below quantify sampling error over test subjects for a *fixed* pair of score vectors, not seed-to-seed variance of pretraining. Read every comparison here as a controlled single-run observation, not an established ranking over retrains (Section 6).

5.1 H1 holds, H2 fails

Table 1 gives the main comparison, Table 9 in Appendix I the full paired inference for all nine precision-matched contrasts, and Figure 2 the trajectories. All six contrasts against the null survive Benjamini–Hochberg correction over that family of nine ($q \leq 0.0299$ throughout). **H1 holds:** constraining the same rectangles to lie on retinal tissue (ENVELOPE) beats unguided masking by +0.0120 at epoch 50 (CI [+0.0068, +0.0173]) and by +0.0062 at epoch 100 (CI [+0.0010, +0.0113]). Avoiding background is worth something, at every epoch we measured, with an interval that excludes zero.

H2 fails. The arms that use the segmentation model most aggressively — shaping targets to trace the retina, or maximising anatomy coverage under a floor — do not improve on the null. At epoch 50, the only epoch at which all five arms have a probe, ANATOMY-V2 gains +0.0013 over RANDOM and COVER gains +0.0002, against +0.0120 for ENVELOPE and +0.0099 for INTENSITY. The policies that merely *place* rectangles on tissue gain roughly an order of magnitude more than the policies that shape targets to match it. Increasing anatomical fidelity past “rectangles on tissue” does not add benefit.

We are careful not to overstate this. At the matched epoch 50 neither ANATOMY-V2 (CI [−0.0055, +0.0082], $p = 0.7153$) nor COVER (CI [−0.0050, +0.0053], $p = 0.9513$) is *worse* than

Table 1: Frozen MeanPool probe, FairVision test ($N=3000$, 1466 positive). All arms continue from one epoch-25 ancestor (AUC 0.8487). Δ is versus RANDOM at the *matched* epoch. The upper block is precision-matched (fp16 throughout) and carries every headline claim. The lower block is fp32. Its epoch-50 Δ is taken against an fp32 null re-probed for that comparison and is matched on epoch and precision. The COVER epoch-75 Δ is marked ‡: no fp32 null exists at that epoch yet, so it is matched on epoch only. The measured fp16/fp32 gap is at most 2×10^{-4} (Appendix J), some forty times below that Δ .

policy	guidance signal	epoch 50		epoch 75		epoch 100	
		AUC	Δ	AUC	Δ	AUC	Δ
RANDOM	none (null)	0.8641	—	0.8723	—	0.8746	—
ENVELOPE	segmenter (location only)	0.8761	+0.0120	0.8803	+0.0080	0.8807	+0.0062
INTENSITY	intensity centroid, no segmenter	0.8740	+0.0099	0.8836	+0.0113	0.8855	+0.0109
ANATOMY-V2	segmenter (target shape)	0.8654	+0.0013	—	—	—	—
COVER $f=.21$	segmenter (coverage)	0.8643	+0.0002	0.8639	−0.0084‡	pending†	pending†




Figure 2: **(a)** Frozen-probe test AUC against pretraining epoch. All arms fork from the same epoch-25 checkpoint (open circle) and differ only in masking policy. Solid lines are the precision-matched fp16 family; dashed lines are fp32 arms and must not be compared vertically against the solid lines. **(b)** The actual inference: paired differences against the unguided null on identical test cases, with 10,000-resample bootstrap intervals. All six intervals exclude zero. We plot paired differences rather than per-arm error bars because marginal intervals carry between-case variance that cancels in a paired comparison, and would understate the evidence.

the null; both are indistinguishable from it, and both are measured against an fp32 null re-probed specifically for this comparison, so they are matched on epoch *and* precision. What the epoch-50 data rule out is not that these policies are harmful, but the expected *monotone improvement with anatomical precision*: their intervals are centred near zero and comfortably contain it, while the two cruder policies produce intervals that exclude it.

Carried further, coverage-constrained masking stalls and falls behind. COVER is the only segmenter-driven policy we continued past epoch 50, and it produces the study’s one *negative* result. Its trajectory is 0.8483, 0.8522, 0.8571, 0.8643, 0.8647, 0.8639: it climbs normally to epoch 50 and then stops. From epoch 50 to 75 it moves -0.0004 ($p = 0.8655$), while the null gains $+0.0082$. By epoch 75 it sits *below* the null by -0.0084 (CI $[-0.0146, -0.0026]$, $p = 0.0058$), below ENVELOPE by -0.0164 and below INTENSITY by -0.0198 . Forcing a fixed fraction of tissue to stay visible does not merely fail to help; it appears to cap what the encoder can still learn once the easy signal is exhausted. Two caveats: these probes are fp32 against fp16 comparators, though the measured precision effect is at most 2×10^{-4} (Appendix J), some forty times below this gap; and only $f=.21$ was pretrained, so this concerns that floor, not a dose-response.

The strongest policy is the one that uses no segmentation model. INTENSITY gives $+0.0109$ over the null at epoch 100 (95% CI $[+0.0057, +0.0160]$, $p < 0.0001$, $q = < 0.0001$), with a margin that is stable across training ($+0.0099, +0.0113, +0.0109$). It exceeds the segmenter-guided ENVELOPE

Table 2: Measured mask geometry, 600 held-out slices, production samplers. “purity” is the fraction of masked cells lying on tissue; “loss slots” is the number of predictor targets contributing to the loss. All values measured. AUC is quoted at the *matched* epoch 50, the only epoch at which all five policies have a probe, so that the geometry-to-AUC association is not read across mixed epochs. All five AUCs are at epoch 50.

policy	anatomy hidden	purity	mask ratio	context kept	loss slots	AUC @ep50
RANDOM	52.2%	31.6%	43.7%	42.1%	157.7	0.8641
INTENSITY	62.2%	41.1%	40.0%	45.6%	158.4	0.8740
ENVELOPE	76.9%	43.5%	46.4%	40.5%	159.9	0.8761
COVER	74.1%	45.3%	43.3%	43.5%	160.0	0.8643
ANATOMY-V2	80.3%	97.3%	21.4%	67.9%	64.0	0.8654




Figure 3: Downstream AUC at matched epoch 50 against mask purity, the fraction of masked patches lying on retinal tissue. The policy whose masks are 97% on-tissue does not separate from the null, while policies at 41–44% purity gain the most. The ordering does not support the intuition that more accurate anatomical masking yields better representations. With four to five arms this is descriptive, not inferential; the same result arranged as a specificity ladder is in Appendix J.

arm at epoch 100 by $+0.0047$ (CI $[+0.0004, +0.0091]$, $p = 0.0294$), though the two are indistinguishable at epochs 50 and 75. Notably, ENVELOPE’s margin over the null *decays* with training ($+0.0120 \rightarrow +0.0080 \rightarrow +0.0062$) while INTENSITY’s does not.

5.2 H3: it is not anatomical targeting

To test H3 we ran each production sampler on 600 held-out slices and measured what it actually does (Table 2, Figure 3). Three observations follow.

First, **anatomical targeting does not order the results**. Over the four rectangle arms, the Spearman correlation between fraction-of-anatomy-hidden and AUC is exactly 0.00: knowing how much anatomy a policy hides tells you nothing about its downstream AUC. Mask *purity* — the share of masked patches lying on tissue — trends negatively (-0.40). Neither is significant at $n=4$ arms, and we present them as descriptive rather than inferential; but neither shows the positive relationship the anatomy-guidance hypothesis predicts. The anatomy arm places 97.3% of masked cells on tissue and does not separate from the null; INTENSITY places 41.1% on tissue and gives the largest gain in the study.

Second, **the anatomy arms are confounded on several axes simultaneously**, and we state this rather than attributing their deficit to shape. Because their targets are compact tissue blobs rather than large rectangles, they hide only 21.4% of the grid against 40–46% for the rectangle arms, leave 67.9% of tokens visible as context against 40–46%, and supply 64 predictor loss slots against 158–160. Any of these could account for their deficit without invoking target shape at all. We therefore do *not* claim that irregular target shape is harmful; that comparison is not identified by this design. Figure 12 makes the four-way divergence explicit.

Third, INTENSITY is the most *spatially consistent* policy we measured. Its per-cell hit map is the most concentrated of any arm (Gini 0.400 versus 0.316–0.338 for the other rectangle arms): it repeatedly asks the encoder to represent the same region of the image. This suggests a mechanism

248 other than anatomical correctness — consistency of the predictive task — which we return to in
249 Section 6.

250 5.3 Subgroup analysis

251 Joining saved test predictions to the released FairVision metadata in deterministic loader order (the
252 join is validated by exact label reconstruction: the label vector rebuilt from the sorted file order
253 matches the stored labels for all 3000 cases in every probe), we audit seven stratifications across 20
254 probes spanning aggregate AUC 0.8483 to 0.8855. Probes that the exclusion rules of Appendix D
255 remove are removed here too.

256 **The ordering of subgroups never changes.** The worst-performing group is the same in every
257 one of the 20 probes for sex (female), race (black), ethnicity (hispanic) and disease severity (mild
258 $(-6, -2]$). No masking policy reorders them. The probes share a test set, an epoch-25 ancestor
259 and a probe seed, and several are different epochs of one branch, so they are *not* independent; we
260 report this as a descriptive regularity and attach no p -value. What it shows is that the disparity is not
261 introduced by any policy we tried.

262 **Every group improves; no gap reliably moves.** A widening max–min gap is easy to misread as
263 harm. It is not. At epoch 100 the strongest policy improves *every* race stratum over the null — black
264 $0.8325 \rightarrow 0.8472 (+0.0147, \text{a larger gain than white's } 0.8741 \rightarrow 0.8853)$ — and the max–min
265 gap widens only because the already-best asian subgroup gains most. Nor is there a defensible gap
266 *trend*. Across checkpoints the race gap correlates with aggregate AUC at $\rho = +0.517$ ($p = 0.0195$,
267 $q = 0.0455$ after Benjamini–Hochberg over the seven attributes tested), but those checkpoints are
268 not independent: they are 20 probes drawn from only 7 pretraining branches, so a branch probed at
269 four epochs contributes four correlated points. Collapsing to one point per branch, the association
270 disappears ($\rho = +0.429$, $p = 0.3374$). We therefore treat the checkpoint-level correlation as pseudo-
271 replicated and make no claim that better models are less fair. The same holds for severity in the
272 opposite direction ($\rho = -0.519$ per checkpoint, -0.750 per branch). Only the sex gap is consistent
273 across both views, and it *narrows* ($q = 0.0039$; branch-level $\rho = -0.857$, $p = 0.0137$). We do not
274 claim that better masking harms any group, and we do not claim it helps equity either.

275 **Early disease improves most; differential benefit is unresolved.** Scoring each severity stratum
276 against the shared pool of all negatives, the strongest policy improves mild disease by $+0.0137$
277 at matched epoch 100, moderate by $+0.0102$, severe by $+0.0063$. These are paired differences
278 on the same subjects and all three intervals exclude zero (mild $[+0.00541, +0.02192]$, severe
279 $[+0.00104, +0.01187]$), so the ordering is resolved rather than a point-estimate artefact. What
280 does not move is the difficulty ordering itself: the severe-to-mild gap stays within 0.0114 of con-
281 stant across all 20 probes while aggregate AUC spans 0.8483–0.8855. For equity the quantity that
282 matters is the paired *difference* between groups' gains; between the worst- and best-served race
283 strata it is -0.00091 (CI $[-0.02379, +0.02139]$) and by sex $+0.00146$ (CI $[-0.00853, +0.01156]$).
284 Both contain zero, so at these subgroup sizes we cannot resolve whether the improvement is evenly
285 distributed, and claim neither. Appendix C gives every stratum. No policy here was designed as a
286 fairness intervention and none should be presented as one.

287 5.4 Numerical precision is not the effect

288 Because the long-horizon arms and the anatomy arms were probed at different autocast settings
289 (Section 4), we measured the precision effect rather than assuming it away. Re-encoding the epoch-
290 100 INTENSITY test split at fp32 reproduces the reported AUC to 9.8×10^{-6} , and a full re-fit of the
291 entire probe pipeline at fp32 shifts every long-horizon arm by less than 2×10^{-4} . Both are two to
292 three orders of magnitude below the effects in Table 1. Appendix J gives the numbers.

293 5.5 Reproducibility

294 The epoch-100 INTENSITY encoder and its trained probe head will be released, together with the
295 stored per-case predictions for every probe in this paper. Links are withheld here for anonymity.
296 Every number in this paper is emitted by a single script from those stored predictions, so tables,
297 figures and prose cannot disagree; that script is released with the code.

6 Discussion and limitations

What the evidence supports. Guidance helps: restricting targets to tissue is worth $+0.0120$ at the matched epoch, and a consistent intensity-located band is worth $+0.0109$ at epoch 100, with paired intervals excluding zero at every epoch measured. A trained segmentation model is not required to obtain this benefit, and using one more aggressively — to shape targets or to maximise anatomy coverage — does not add to it and can subtract. The practical reading, bounded by the caveats below, is that for *this* task the segmentation stage bought nothing we could measure, and a one-line intensity statistic recovered the whole benefit. We would not generalise that to tasks where anatomy is the output, such as layer segmentation or lesion localisation.

What it does not support. We cannot claim anatomy-shaped masking is *harmful*: at the matched epoch it sits $+0.0013$ from the null with an interval $([-0.0055, +0.0082])$ that comfortably contains zero, so it is indistinguishable rather than worse. Nor can we claim target *shape* is the operative variable, for the reasons in Table 2. Nor can we claim the mechanism behind INTENSITY’s advantage. Two explanations remain live: consistency of location (the encoder is repeatedly forced to represent the same region), and masking ratio or task difficulty. Distinguishing them requires an arm that randomises the band’s vertical position while holding area and context fixed, which we did not run.

One pretraining run per policy. The paired bootstrap quantifies sampling error on a fixed test set; it does not quantify seed-to-seed variance of pretraining. With $n=1$ continuation per arm the *ranking* is not statistically established even where individual pairwise differences are significant on this test set. Bouthillier et al. [2021] show that single-seed comparisons routinely mistake optimisation noise for method effects. We can bound one component: with five probe seeds on two fixed encoders the probe-seed standard deviation is 0.0003 and 0.0018, and the arms stay fully separated across all five, so probe noise is an order of magnitude below the effects in Table 1. Those are *technical* replicates and do not estimate pretraining noise, which would need at least three paired continuations from the shared checkpoint with the continuation as the unit of replication. That is the single most valuable follow-up and we do not claim its result in advance.

Scope, coverage and an incomplete arm. All results are glaucoma classification from FairVision OCT; generality to other modalities, diseases, or segmentation-quality regimes is untested, and the anatomy arms depend on MIRAGE’s prediction quality on this data, so a better segmenter might change the conclusion. Only $f=0.21$ was pretrained for COVER, so we report no dose-response over the visible-tissue floor and make no claim that any floor is optimal. That arm also has probes only at epochs 27, 30, 34, 50 and 73: it was deliberately halted before epoch 100, its later probes are marked pending rather than reported, and we draw no conclusion about its endpoint. Its epoch-50 value is included because it is a matched-epoch comparison against every other arm. Finally, one dataset and one test split were reused throughout this programme’s history, and policies, checkpoints and analyses were chosen after repeated inspection of that split; the intervals we report are therefore best read as descriptive rather than as confirmatory inference.

Broader impact and ethics. This study uses a public, de-identified retinal dataset under its release terms and involves no new data collection or human subjects. No model reported here is clinically validated or intended for deployment. Appendix G gives the full statement.

7 Conclusion

We tested the intuition that masked predictive pretraining in medical imaging should target clinically meaningful anatomy. Under matched compute, schedule, and a shared ancestor checkpoint, guidance helped — but anatomical precision did not. Restricting targets to retinal tissue improved AUC over unguided masking by $+0.0120$ at the matched epoch, while shaping targets to trace the segmented retina, or maximising anatomy coverage, left performance indistinguishable from not guiding at all; and the best policy consulted no segmentation model. Mask geometry is uncorrelated with how much anatomy is hidden, and a subgroup audit over 20 probes shows every group and every disease stage improving — mild disease most of all — while the difficulty ordering between them is untouched. We suggest the operative variable is the consistency of the predictive task rather than its anatomical correctness, and will release weights and the reproduction protocol.

References

- Mahmoud Assran, Quentin Duval, Ishan Misra, Piotr Bojanowski, Pascal Vincent, Michael Rabbat, Yann LeCun, and Nicolas Ballas. Self-supervised learning from images with a joint-embedding predictive architecture. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 15619–15629, 2023. doi: 10.1109/CVPR52729.2023.01499.
- Shekoofeh Azizi, Basil Mustafa, Fiona Ryan, Zachary Beaver, Jan Freyberg, Jonathan Deaton, Aaron Loh, Alan Karthikesalingam, Simon Kornblith, Ting Chen, et al. Big self-supervised models advance medical image classification. In *Proceedings of the IEEE/CVF International Conference on Computer Vision*, 2021.
- Alexei Baevski, Wei-Ning Hsu, Qiantong Xu, Arun Babu, Jiatao Gu, and Michael Auli. data2vec: A general framework for self-supervised learning in speech, vision and language. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 1298–1312. PMLR, 2022. URL <https://proceedings.mlr.press/v162/baevski22a.html>.
- Randall Balestriero and Yann LeCun. LeJEPa: Provable and scalable self-supervised learning without the heuristics, 2025. URL <https://arxiv.org/abs/2511.08544>.
- Hangbo Bao, Li Dong, Songhao Piao, and Furu Wei. BEiT: BERT pre-training of image transformers. In *International Conference on Learning Representations*, 2022. URL <https://openreview.net/forum?id=p-BhZSz59o4>.
- Adrien Bardes, Quentin Garrido, Jean Ponce, Xinlei Chen, Michael Rabbat, Yann LeCun, Mahmoud Assran, and Nicolas Ballas. Revisiting feature prediction for learning visual representations from video. *Transactions on Machine Learning Research*, 2024. URL <https://openreview.net/forum?id=QaCCuDFbk2>.
- Xavier Bouthillier, Pierre Delaunay, Mirko Bronzi, Assya Trofimov, Brennan Nichyporuk, Justin Szeto, Nazanin Mohammadi Sepahvand, Edward Raff, Kanika Madan, Vikram Voleti, Samira Ebrahimi Kahou, Vincent Michalski, Tal Arbel, Chris Pal, Gael Varoquaux, and Pascal Vincent. Accounting for variance in machine learning benchmarks. *Proceedings of Machine Learning and Systems*, 3:747–769, 2021. URL https://proceedings.mlsys.org/paper_files/paper/2021/hash/0184b0cd3cfb185989f858a1d9f5c1eb-Abstract.html.
- Alberto M. Ceballos Arroyo, Jisoo Kim, Chu-Hsuan Lin, Lei Qin, Geoffrey S. Young, and Huaizu Jiang. Anatomically-guided masked autoencoder pre-training for aneurysm detection. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 5693–5702, 2026. doi: 10.1109/WACV61042.2026.00552.
- Haijian Chen, Wendong Zhang, Yunbo Wang, and Xiaokang Yang. Improving masked autoencoders by learning where to mask. In *Pattern Recognition and Computer Vision*, pages 377–390. Springer, 2023a. doi: 10.1007/978-981-99-8543-2_31.
- Zekai Chen, Devansh Agarwal, Kshitij Aggarwal, Wiem Safta, Mariann Micsinai Balan, and Kevin Brown. Masked image modeling advances 3d medical image analysis. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 1970–1980, 2023b. doi: 10.1109/WACV56688.2023.00201.
- Elizabeth R. DeLong, David M. DeLong, and Daniel L. Clarke-Pearson. Comparing the areas under two or more correlated receiver operating characteristic curves: A nonparametric approach. *Biometrics*, 44(3):837–845, 1988. doi: 10.2307/2531595. URL <https://doi.org/10.2307/2531595>.
- Judy Wawira Gichoya, Imon Banerjee, Ananth Reddy Bhimireddy, John L. Burns, Leo Anthony Celi, Li-Ching Chen, Ramon Correa, Natalie Dullerud, Marzyeh Ghassemi, Shih-Cheng Huang, Po-Chih Kuo, Matthew P. Lungren, Lyle J. Palmer, Brandon J. Price, Saptarshi Purkayastha, Ayis T. Pyrros, Lauren Oakden-Rayner, Chima Okechukwu, Laleh Seyyed-Kalantari, Hari Trivedi, Ryan Wang, Zachary Zaiman, and Haoran Zhang. AI recognition of patient race in medical imaging: A modelling study. *The Lancet Digital Health*, 4(6):e406–e414, 2022. doi: 10.1016/S2589-7500(22)00063-2.

400 Han Guo, Ramtin Hosseini, Ruiyi Zhang, Sai Ashish Somayajula, Ranak Roy Chowdhury, Rajesh K.
401 Gupta, and Pengtao Xie. Downstream task guided masking learning in masked autoencoders using
402 multi-level optimization. *Transactions on Machine Learning Research*, 2025. arXiv:2402.18128.

403 Kaiming He, Xinlei Chen, Saining Xie, Yanghao Li, Piotr Dollár, and Ross Girshick. Masked autoen-
404 coders are scalable vision learners. In *Proceedings of the IEEE/CVF Conference on Computer Vi-*
405 *sion and Pattern Recognition*, pages 16000–16009, 2022. doi: 10.1109/CVPR52688.2022.01553.

406 Xiangteng He, Shunsuke Sakai, Shivam Chandhok, Sara Beery, Kun Yuan, Nicolas Padoy, Tatsuhito
407 Hasegawa, and Leonid Sigal. DSeq-JEPA: Discriminative sequential joint-embedding predictive
408 architecture, 2025. URL <https://arxiv.org/abs/2511.17354>. Accepted at ECCV 2026;
409 final proceedings metadata unavailable at verification time.

410 De-Xing Huang, Xiao-Hu Zhou, Mei-Jiang Gui, Xiao-Liang Xie, Shi-Qi Liu, Shuang-Yi Wang,
411 Tian-Yu Xiang, Rui-Ze Ma, Nu-Fang Xiao, and Zeng-Guang Hou. VasoMIM: Vascular anatomy-
412 aware masked image modeling for vessel segmentation. In *Proceedings of the AAAI Conference*
413 *on Artificial Intelligence*, volume 40, pages 4985–4993, 2026. doi: 10.1609/aaai.v40i6.42503.

414 Yijin Huang, Junyan Lyu, Pujin Cheng, Roger Tam, and Xiaoying Tang. SSiT: Saliency-guided
415 self-supervised image transformer for diabetic retinopathy grading. *IEEE Journal of Biomedical*
416 *and Health Informatics*, 28(5):2806–2817, 2024. doi: 10.1109/JBHI.2024.3362878. URL <https://doi.org/10.1109/JBHI.2024.3362878>.
417 [//doi.org/10.1109/JBHI.2024.3362878](https://doi.org/10.1109/JBHI.2024.3362878).

418 Ruinan Jin, Zikang Xu, Yuan Zhong, Qingsong Yao, Qi Dou, S. Kevin Zhou, and Xiaoxiao
419 Li. FairMedFM: Fairness benchmarking for medical imaging foundation models. In *Ad-*
420 *vances in Neural Information Processing Systems*, volume 37, pages 111318–111357, 2024.
421 doi: 10.52202/079017-3535. URL [https://proceedings.neurips.cc/paper_files/](https://proceedings.neurips.cc/paper_files/paper/2024/hash/c9826b9ea5e1b49b256329934a578d83-Abstract-Datasets_and_Benchmarks_Track.html)
422 [paper/2024/hash/c9826b9ea5e1b49b256329934a578d83-Abstract-Datasets_and_](https://proceedings.neurips.cc/paper_files/paper/2024/hash/c9826b9ea5e1b49b256329934a578d83-Abstract-Datasets_and_Benchmarks_Track.html)
423 [Benchmarks_Track.html](https://proceedings.neurips.cc/paper_files/paper/2024/hash/c9826b9ea5e1b49b256329934a578d83-Abstract-Datasets_and_Benchmarks_Track.html).

424 Ioannis Kakogeorgiou, Spyros Gidaris, Bill Psomas, Yannis Avrithis, Andrei Bursuc, Konstantinos
425 Karantzas, and Nikos Komodakis. What to hide from your students: Attention-guided masked
426 image modeling. In *Computer Vision – ECCV 2022*, pages 300–318. Springer, 2022. doi: 10.
427 1007/978-3-031-20056-4_18.

428 Agostina J. Larrazabal, Nicolás Nieto, Victoria Peterson, Diego H. Milone, and Enzo Ferrante.
429 Gender imbalance in medical imaging datasets produces biased classifiers for computer-aided
430 diagnosis. *Proceedings of the National Academy of Sciences*, 117(23):12592–12594, 2020. doi:
431 10.1073/pnas.1919012117.

432 Jin Lee, Vu Dang, Gwang-Hyun Yu, Anh Le, Zahid Rahman, Jin-Ho Jang, Heonzoo Lee, Kun-Yung
433 Kim, Jin-Sul Kim, and Jin-Young Kim. HU-based foreground masking for 3D medical masked
434 image modeling. In *Medical Image Computing and Computer Assisted Intervention (MICCAI)*,
435 2025. doi: 10.1007/978-3-032-09569-5_13. arXiv:2509.07534.

436 Gang Li, Heliang Zheng, Daqing Liu, Chaoyue Wang, Bing Su, and Changwen Zheng.
437 SemMAE: Semantic-guided masking for learning masked autoencoders. In *Advances in*
438 *Neural Information Processing Systems*, volume 35, pages 14290–14302, 2022. doi:
439 10.52202/068431-1039. URL [https://proceedings.neurips.cc/paper_files/paper/](https://proceedings.neurips.cc/paper_files/paper/2022/hash/5db4e0a62cb5b4a5a70f1fdc9a78a4cc-Abstract-Conference.html)
440 [2022/hash/5db4e0a62cb5b4a5a70f1fdc9a78a4cc-Abstract-Conference.html](https://proceedings.neurips.cc/paper_files/paper/2022/hash/5db4e0a62cb5b4a5a70f1fdc9a78a4cc-Abstract-Conference.html).

441 Yuheng Li, Tianyu Luan, Yizhou Wu, Shaoyan Pan, Yenho Chen, and Xiaofeng Yang. AnatoMask:
442 Enhancing medical image segmentation with reconstruction-guided self-masking. In *Computer*
443 *Vision – ECCV 2024*, pages 146–163. Springer, 2024. doi: 10.1007/978-3-031-73027-6_9.

444 Zixuan Liu, Hanwen Xu, Addie Woicik, Linda G. Shapiro, Marian Blazes, Yue Wu, Verena Steffen,
445 Catherine A. Cukras, Cecilia S. Lee, Miao Zhang, Aaron Y. Lee, and Sheng Wang. A three-
446 dimensional multi-modal foundation model for optical coherence tomography. *Nature Biomedical*
447 *Engineering*, 2026. doi: 10.1038/s41551-026-01662-2.

448 Yan Luo, Muhammad Osama Khan, Yu Tian, Min Shi, Zehao Dou, Tobias Elze, Yi Fang, and
449 Mengyu Wang. FairVision: Equitable deep learning for eye disease screening via fair identity
450 scaling, 2024a. URL <https://arxiv.org/abs/2310.02492>.

451 Yan Luo, Yu Tian, Min Shi, Louis R. Pasquale, Lucy Q. Shen, Nazlee Zebardast, Tobias Elze, and
452 Mengyu Wang. Harvard glaucoma fairness: A retinal nerve disease dataset for fairness learning
453 and fair identity normalization. *IEEE Transactions on Medical Imaging*, 43(7):2623–2633, 2024b.
454 doi: 10.1109/TMI.2024.3377552.

455 Stefan Maetschke, Bhavna Antony, Hiroshi Ishikawa, Gadi Wollstein, Joel Schuman, and Rahil
456 Garnavi. A feature agnostic approach for glaucoma detection in OCT volumes. *PLOS ONE*, 14
457 (7), 2019.

458 Lena Maier-Hein, Annika Reinke, Patrick Godau, Minu D. Tizabi, Florian Buettner, Evangelia
459 Christodoulou, Ben Glocker, Fabian Isensee, Jens Kleesiek, Michal Kozubek, Mauricio Reyes,
460 Michael A. Riegler, Manuel Wiesenfarth, A. Emre Kavur, Carole H. Sudre, Michael Baumgartner,
461 Matthias Eisenmann, Doreen Heckmann-Nötzel, Tim Radsch, Laura Acion, Michela Antonelli,
462 Tal Arbel, Spyridon Bakas, Arriel Benis, Matthew B. Blaschko, M. Jorge Cardoso, Veronika Chep-
463 lygina, Beth A. Cimini, Gary S. Collins, Keyvan Farahani, Luciana Ferrer, Adrian Galdran, Bram
464 van Ginneken, Robert Haase, Daniel A. Hashimoto, Michael M. Hoffman, Merel Huisman, Pierre
465 Jannin, Charles E. Kahn, Dagmar Kainmueller, Bernhard Kainz, Alexandros Karargyris, Alan
466 Karthikesalingam, Florian Kofler, Annette Kopp-Schneider, Anna Kreshuk, Tahsin Kurc, Ben-
467 nett A. Landman, Geert Litjens, Amin Madani, Klaus Maier-Hein, Anne L. Martel, Peter Matt-
468 son, Erik Meijering, Bjoern Menze, Karel G. M. Moons, Henning Müller, Brennan Nihyporuk,
469 Felix Nickel, Jens Petersen, Nasir Rajpoot, Nicola Rieke, Julio Saez-Rodriguez, Clara I. Sánchez,
470 Shravya Shetty, Maarten van Smeden, Ronald M. Summers, Abdel A. Taha, Aleksei Tiulpin,
471 Sotirios A. Tsaftaris, Ben Van Calster, Gaël Varoquaux, and Paul F. Jäger. Metrics reloaded:
472 Recommendations for image analysis validation. *Nature Methods*, 21(2):195–212, 2024. doi:
473 10.1038/s41592-023-02151-z. URL <https://doi.org/10.1038/s41592-023-02151-z>.

474 Jiawei Mao, Shujian Guo, Yuanqi Chang, Xuesong Yin, and Binling Nie. Medical supervised
475 masked autoencoders: Crafting a better masking strategy and efficient fine-tuning schedule for
476 medical image classification, 2023. URL <https://arxiv.org/abs/2305.05871>.

477 José Morano, Botond Fazekas, Emese Sükei, Ronald Fecso, Taha Emre, Markus Gumpinger, Georg
478 Faustmann, Marzieh Oghbaie, Ursula Schmidt-Erfurth, and Hrvoje Bogunović. Multimodal
479 foundation model and benchmark for comprehensive retinal OCT image analysis. *npj Digital*
480 *Medicine*, 8:576, 2025. doi: 10.1038/s41746-025-01852-3.

481 Erfan Noury, Suria S. Mannil, Robert T. Chang, An Ran Ran, Carol Y. Cheung, Suman S. Thapa,
482 Harsha L. Rao, Srilakshmi Dasari, Mohammed Riyazuddin, Dolly Chang, Sriharsha Nagaraj,
483 Clement C. Tham, and Reza Zadeh. Deep learning for glaucoma detection and identification of
484 novel diagnostic areas in diverse real-world datasets. *Translational Vision Science & Technology*,
485 11(5):11, 2022. doi: 10.1167/tvst.11.5.11.

486 Nancy A. Obuchowski. Nonparametric analysis of clustered ROC curve data. *Biometrics*, 53(2):
487 567–578, 1997. doi: 10.2307/2533958. URL <https://doi.org/10.2307/2533958>.

488 Theodoros Pissas, Pablo Márquez-Neila, Sebastian Wolf, Martin S. Zinkernagel, and Raphael Sznit-
489 man. Masked image modelling for retinal OCT understanding. In *Ophthalmic Medical Image*
490 *Analysis*, pages 115–125. Springer, 2024. doi: 10.1007/978-3-031-73119-8_12.

491 Jianing Qiu, Jian Wu, Hao Wei, Peilun Shi, Mingqing Zhang, Yunyun Sun, Lin Li, Hanruo Liu,
492 Hongyi Liu, Simeng Hou, Yuyang Zhao, Xuehui Shi, Junfang Xian, Xiaoxia Qu, Sirui Zhu, Lijie
493 Pan, Xiaoniao Chen, Xiaojia Zhang, Shuai Jiang, Kebin Wang, Chenlong Yang, Mingqiang
494 Chen, Sujie Fan, Jianhua Hu, Aiguo Lv, Hui Miao, Li Guo, Shujun Zhang, Cheng Pei, Xiaojuan
495 Fan, Jianqin Lei, Ting Wei, Junguo Duan, Chun Liu, Xiaobo Xia, Siqi Xiong, Junhong Li, Kyle
496 Lam, Benny Lo, Yih Chung Tham, Tien Yin Wong, Ningli Wang, and Wu Yuan. Development and
497 validation of a multimodal multitask vision foundation model for generalist ophthalmic artificial
498 intelligence. *NEJM AI*, 1(12), 2024. doi: 10.1056/AIoa2300221.

499 An Ran Ran, Carol Y. Cheung, Xi Wang, Hao Chen, Lu-Yang Luo, Poemen P. Chan, Mandy O. M.
500 Wong, Robert T. Chang, Suria S. Mannil, Alvin L. Young, Hon-Wah Yung, Chi Pui Pang, Pheng-
501 Ann Heng, and Clement C. Tham. Detection of glaucomatous optic neuropathy with spectral-
502 domain optical coherence tomography: A retrospective training and validation deep-learning anal-
503 ysis. *The Lancet Digital Health*, 1(4):e172–e182, 2019. doi: 10.1016/S2589-7500(19)30085-8.

Carolyn M. Rutter. Bootstrap estimation of diagnostic accuracy with patient-clustered data. *Academic Radiology*, 7(6):413–419, 2000. doi: 10.1016/S1076-6332(00)80381-5. URL [https://doi.org/10.1016/S1076-6332\(00\)80381-5](https://doi.org/10.1016/S1076-6332(00)80381-5).

Laleh Seyyed-Kalantari, Haoran Zhang, Matthew B. A. McDermott, Irene Y. Chen, and Marzyeh Ghassemi. Underdiagnosis bias of artificial intelligence algorithms applied to chest radiographs in under-served patient populations. *Nature Medicine*, 27(12):2176–2182, 2021. doi: 10.1038/s41591-021-01595-0.

Min Shi, Yan Luo, Yu Tian, Lucy Q. Shen, Tobias Elze, Nazlee Zebardast, Mohammad Eslami, Saber Kazeminasab, Michael V. Boland, David S. Friedman, Louis R. Pasquale, and Mengyu Wang. Equitable artificial intelligence for glaucoma screening with fair identity normalization. *npj Digital Medicine*, 8(1), 2025. doi: 10.1038/s41746-025-01432-5.

Yuge Shi, N. Siddharth, Philip H. S. Torr, and Adam R. Kosiorek. Adversarial masking for self-supervised learning. In *Proceedings of the 39th International Conference on Machine Learning*, volume 162 of *Proceedings of Machine Learning Research*, pages 20026–20040. PMLR, 2022. URL <https://proceedings.mlr.press/v162/shi22d.html>.

Leon Sick, Dominik Engel, Pedro Hermosilla, and Timo Ropinski. Attention-guided masked autoencoders for learning image representations. In *Proceedings of the IEEE/CVF Winter Conference on Applications of Computer Vision*, pages 836–846, 2025. doi: 10.1109/WACV61041.2025.00091. URL https://openaccess.thecvf.com/content/WACV2025/html/Sick_Attention-Guided_Masked_Autoencoders_for_Learning_Image_Representations_WACV_2025_paper.html.

Aiham Taleb, Winfried Loetzsch, Noel Danz, Julius Severin, Thomas Gärtner, Benjamin Bergner, and Christoph Lippert. 3d self-supervised methods for medical imaging. In *Advances in Neural Information Processing Systems*, volume 33, pages 18158–18172, 2020. URL <https://proceedings.neurips.cc/paper/2020/hash/d2dc6368837861b42020ee72b0896182-Abstract.html>.

Haochen Wang, Kaiyou Song, Junsong Fan, Yuxi Wang, Jin Xie, and Zhaoxiang Zhang. Hard patches mining for masked image modeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10375–10385, 2023a. doi: 10.1109/CVPR52729.2023.01000.

Haochen Wang, Kaiyou Song, Junsong Fan, Yuxi Wang, Jin Xie, and Zhaoxiang Zhang. Hard patches mining for masked image modeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 10375–10385, 2023b. doi: 10.1109/CVPR52729.2023.01000. URL https://openaccess.thecvf.com/content/CVPR2023/html/Wang_Hard_Patches_Mining_for_Masked_Image_Modeling_CVPR_2023_paper.html.

Ruilang Wang, Shuotong Xu, Bowen Liu, Runlin Huang, Donglong Chen, and Weifeng Su. Mask what matters: Controllable text-guided masking for self-supervised medical image analysis, 2025. URL <https://arxiv.org/abs/2509.23054>.

Zhenda Xie, Zheng Zhang, Yue Cao, Yutong Lin, Jianmin Bao, Zhuliang Yao, Qi Dai, and Han Hu. SimMIM: A simple framework for masked image modeling. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 9653–9663, 2022. doi: 10.1109/CVPR52688.2022.00943.

Kai Yu, Yang Zhou, Yang Bai, Zhi Da Soh, Xinxing Xu, Rick Siow Mong Goh, Ching-Yu Cheng, and Yong Liu. UrFound: Towards universal retinal foundation models via knowledge-guided masked modeling. In *Medical Image Computing and Computer Assisted Intervention – MICCAI 2024*, pages 753–762. Springer, 2024. doi: 10.1007/978-3-031-72390-2_70.

Jinghao Zhou, Chen Wei, Huiyu Wang, Wei Shen, Cihang Xie, Alan Yuille, and Tao Kong. iBOT: Image BERT pre-training with online tokenizer. In *International Conference on Learning Representations*, 2022.

554 Yukun Zhou, Mark A. Chia, Siegfried K. Wagner, Murat Seçkin Ayhan, Dominic J. Williamson,
555 Robbert R. Struyven, Timing Liu, Moucheng Xu, Mateo G. Lozano, Peter Woodward-Court, et al.
556 A foundation model for generalizable disease detection from retinal images. *Nature*, 622:156–163,
557 2023. doi: 10.1038/s41586-023-06555-x.

558 A All frozen probes

559 Table 3 lists every frozen MeanPool probe in the study.

Table 3: Every frozen probe in the study, including runs excluded from the analysis and the reasons. Protocol is identical throughout except where the precision column says otherwise: FairVision test $N=3000$, seed 42, frozen encoder, MeanPool + LayerNorm + Linear. “Precision” is the probe-time numerical precision, inferred from the stored prediction dtype, and is distinct from the pretraining target precision discussed in Appendix D.

policy	epoch	precision	test AUC	95% CI
ANATOMY-V1	30	fp32	0.8583	[0.8450, 0.8712]
ANATOMY-V2	35	fp32	0.8661	[0.8531, 0.8787]
ANATOMY-V2	40	fp32	0.8683	[0.8553, 0.8807]
ANATOMY-V2	50	fp32	0.8654	[0.8522, 0.8781]
ancestor	25	fp32	0.8487	[0.8349, 0.8621]
COVER	27	fp32	0.8483	[0.8346, 0.8617]
COVER	30	fp32	0.8522	[0.8386, 0.8654]
COVER	34	fp32	0.8571	[0.8437, 0.8700]
COVER	50	fp32	0.8643	[0.8513, 0.8768]
COVER	73	fp32	0.8647	[0.8517, 0.8772]
COVER	75	fp32	0.8639	[0.8507, 0.8763]
ENVELOPE	30	fp32	0.8539	[0.8405, 0.8669]
ENVELOPE	50	fp16	0.8761	[0.8635, 0.8879]
ENVELOPE	50	fp32	0.8761	[0.8635, 0.8879]
ENVELOPE	75	fp16	0.8803	[0.8677, 0.8922]
ENVELOPE	75	fp32	0.8803	[0.8677, 0.8922]
ENVELOPE	100	fp16	0.8807	[0.8683, 0.8926]
ENVELOPE	100	fp32	0.8808	[0.8683, 0.8927]
INTENSITY	50	fp16	0.8740	[0.8615, 0.8862]
INTENSITY	50	fp32	0.8740	[0.8614, 0.8862]
INTENSITY	75	fp16	0.8836	[0.8714, 0.8955]
INTENSITY	100	fp16	0.8855	[0.8735, 0.8971]
INTENSITY	100	fp32	0.8853	[0.8732, 0.8971]
RANDOM	50	fp16	0.8641	[0.8510, 0.8769]
RANDOM	50	fp32	0.8641	[0.8511, 0.8769]
RANDOM	75	fp16	0.8723	[0.8593, 0.8849]
RANDOM	100	fp16	0.8746	[0.8618, 0.8869]
RANDOM	100	fp32	0.8745	[0.8617, 0.8868]
anatomy-v2	75	fp32	0.8625	<i>excluded</i>
anatomy-v2	92	fp32	0.8604	<i>excluded</i>
random	30	fp32	0.8558	<i>retracted</i>
random	50	fp32	0.8590	<i>retracted</i>
random	75	fp32	0.8612	<i>retracted</i>
random	100	fp32	0.8607	<i>retracted</i>

560 B Masking policies across the volume

561 Figure 1 shows one B-scan for legibility. Figure 4 repeats the same rendering on five B-scans
562 sampled across the depth of a volume, so that the stability of each policy under changing retinal
563 geometry can be inspected. The samplers, the colour assignment, and the anatomy contour are
564 identical to Figure 1.

C Subgroup and severity analysis

All numbers below come from saved test predictions joined to FairVision’s released metadata in deterministic loader order; the join is validated by exact label reconstruction (3000/3000 labels recovered in every probe). The 20 probes used here are exactly those the inventory marks valid: the four retracted coverage probes of Appendix D and the two precision-spliced ANATOMY-V2 probes are excluded here as everywhere else in the paper.

Why we report two sample sizes. The 20 probes span only 7 pretraining branches, because several are different epochs of one training run. Probes within a branch share an encoder lineage, a seed and the test split, so treating them as 20 independent observations is pseudo-replication. Table 4 therefore reports each trend twice: once per checkpoint, and once after averaging within each branch. Where the two disagree, the branch-level figure is the one to believe.

Table 4: Subgroup gap trends. “gap range” is the max–min subgroup AUC gap observed across the 20 valid probes. q is Benjamini–Hochberg across the seven attributes tested, which is necessary because seven trends were examined. Only the sex gap and the severity gap survive correction, and both *narrow* as aggregate AUC rises.

attribute	worst group	gap range	per checkpoint ($n=20$)			per branch ($n=7$)	
			ρ	p	q	ρ	p
sex	female (20/20)	0.0272–0.0396	-0.702	0.0006	0.0039	-0.857	0.0137
race	black (20/20)	0.0433–0.0935	+0.517	0.0195	0.0455	+0.429	0.3374
ethnicity	hispanic (20/20)	0.0203–0.0661	+0.355	0.1247	0.1746	+0.464	0.2939
language	other languages (16/20)	0.0518–0.0805	+0.254	0.2796	0.2796	+0.571	0.1802
marital status	single (20/20)	0.0349–0.0654	-0.383	0.0951	0.1665	-0.750	0.0522
age	<60 (19/20)	0.0474–0.0704	+0.304	0.1929	0.2251	+0.536	0.2152
disease severity	mild ($-6, -2$] (20/20)	0.1286–0.1400	-0.519	0.0191	0.0455	-0.750	0.0522

Race. The black subgroup has the lowest AUC in every one of the 20 probes ($n=431$ black, $n=251$ asian, $n=2318$ white). At matched epoch 100 the black subgroup improves from 0.8325 under RANDOM to 0.8472 under INTENSITY (+0.0147), a larger absolute gain than the white subgroup’s (0.8741 $\rightarrow$ 0.8853). The max–min gap nonetheless widens slightly, because the asian subgroup — already the highest — gains most (0.9033 $\rightarrow$ 0.9189). A widening max–min gap with universally rising subgroup AUCs is not evidence of harm, and we do not present it as such. The checkpoint-level correlation between aggregate AUC and racial gap is $\rho = +0.517$ ($p=0.0195$), which fails BH correction ($q=0.0455$) and vanishes at branch level ($\rho = +0.429$, $p=0.3374$). We therefore explicitly do *not* claim that more accurate models are less fair.

Disease severity. FairVision’s label is defined by visual-field mean deviation (md): every test volume with $md \leq -2$ is positive. Severity therefore cannot be used as an ordinary subgroup axis, because within-bin AUC is undefined — a trap that silently yields nothing under naive stratification. Instead each positive stratum is scored against the shared pool of all 1534 negatives.

Table 5: Severity-stratified AUC at matched epoch 100. Each positive stratum is scored against the shared pool of all 1534 negatives. Every stratum improves, and mild disease improves most.

stratum	n_+	RANDOM	ENVELOPE	INTENSITY	Δ
severe ($md \leq -12$)	334	0.9525	0.9559	0.9588	+0.0063
moderate ($-12 < md \leq -6$)	460	0.9030	0.9103	0.9131	+0.0102
mild ($-6 < md \leq -2$)	672	0.8164	0.8232	0.8301	+0.0137

The severe-to-mild gap is 0.1361 for the null and 0.1286 for the strongest policy, and stays within 0.0114 (0.1286–0.1400) across all 20 probes while aggregate AUC spans 0.8483–0.8855. Target placement raises the whole severity curve — mild disease most of all — without altering the ordering that makes early disease hard. We note that mild functional loss is harder to separate from normal OCT by construction, so this ordering is expected clinically; what the data add is that no masking policy we tested changes it.

Table 6: Paired change in subgroup AUC from RANDOM to INTENSITY at matched epoch 100. Each bootstrap draw resamples subjects once and applies the same resampled set to both arms and every stratum, so these are paired differences rather than differences of independent estimates.

stratum	n	Δ AUC	95% CI	excludes 0
severity, mild	2206	+0.01372	[+0.00541, +0.02192]	yes
severity, moderate	1994	+0.01016	[+0.00262, +0.01757]	yes
severity, severe	1868	+0.00626	[+0.00104, +0.01187]	yes
race, White	2318	+0.01129	[+0.00518, +0.01731]	yes
race, Black	431	+0.01467	[-0.00055, +0.03060]	no
race, Asian	251	+0.01558	[-0.00079, +0.03283]	no
sex, Female	1716	+0.01152	[+0.00412, +0.01890]	yes
sex, Male	1284	+0.01006	[+0.00335, +0.01678]	yes

Limitations. One dataset, one test split, reused throughout this programme’s history. Because policies, checkpoints and analyses were chosen after repeated inspection of that split, the p -values here should be read as descriptive rather than as confirmatory inference. Race and ethnicity are self-reported categories inherited from FairVision’s coding; at these subgroup sizes we cannot speak to intersectional groups. All figures are frozen mean-pool linear probes with AUC as the only metric; we report no subgroup calibration, no sensitivity at a fixed specificity, and no threshold-transfer analysis, so this audit should not be read as a complete fairness evaluation. Disparities could differ under fine-tuning or a stronger head.

D Excluded runs

Four probes of an earlier coverage arm are excluded from all analysis. They were trained with half-precision EMA targets and a different encoder-truncation policy from the arms they would be compared against, so their deficit is not attributable to masking. We report their existence for completeness and do not use them to support any claim. Similarly, the ANATOMY-V2 probes at epochs 75 and 92 follow a change of target precision partway through that arm’s training; they are listed in Table 3 but excluded from the matched-epoch comparisons in Table 1.

E Occlusion attribution: where the classifier looks

The analyses in this appendix were run on the three *fine-tuned* probes (mean-pool, cross-attention pool, and a depth-1 attentive probe) that tie at test AUC ≈ 0.887 , not on the frozen probes used for the arm comparison in the main text. They characterise what the downstream classifier uses, and they are reported here because two of the findings are cautionary rather than positive.

Method. Attribution is by *occlusion*, which is architecture-agnostic and causal. Gradient $\times$ input is not used because two of the three probes are non-linear in the slice-token matrix, where gradient attribution is misleading. For slice-level attribution each slice token is zeroed in turn and Δ logit recorded; for window-level, seven consecutive slices; for patch-level, leave-one-out over the 256 patch tokens of a target slice.

A bimodal attribution curve that is an artefact, not anatomy. The population-averaged curve peaks at native slice ≈ 63 and ≈ 137 with a dip near 95. The tempting reading — that the model has discovered the superior and inferior optic disc rim — is **wrong**, and we report it because it is exactly the kind of claim a plausible-looking attribution figure invites. Three tests reject it. First, if the two peaks were bilateral anatomy then within a diseased eye they should co-occur, giving positive per-volume correlation between them; the observed correlation is slightly *negative* (-0.22 , -0.07 , -0.14 across the three probes). Second, clustering the per-volume curves into two groups yields clusters that are near-perfect *mirror images* along the slice axis ($\text{corr}(c_1, \text{flip}(c_2)) = 0.971$ and 0.988 for the two well-structured probes, against raw correlations of -0.124 and -0.478). Third, using that cluster label as pseudo-laterality and flipping one cluster collapses the bimodality asymmetrically. The signature is OD/OS axial storage: each eye carries its dominant signal on one side

of the disc, right and left eyes are stored with opposite axial orientation, and population averaging manufactures an illusion of symmetry.

Errors are weaker signal, not different anatomy. Stratifying by outcome, false-negative curves are scaled-down true-positive curves with the same shape, and false positives likewise mirror true negatives. Attribution structure is also near-invariant to prediction confidence ($|r| \leq 0.25$ between peak contribution magnitude and $|\text{logit}|$). The error rate at threshold 0.5 therefore reflects signal-strength saturation on hard cases, not the model attending to the wrong structures.

Why this is in a masking paper. Two of these results bear directly on the main claim. The classifier’s evidence is spread across the middle two-thirds of the volume and across most patches within a slice, rather than concentrated on a compact anatomical target. A masking policy that sharpens targets onto a precise anatomical boundary is therefore optimising for a spatial prior the downstream task does not appear to need, which is consistent with Section 5.2 finding no relationship between anatomical specificity and downstream AUC. The OD/OS result is a methodological caution of the same family as our precision and epoch-matching checks: a figure that looks anatomically meaningful can be an artefact of data storage, and we would rather report that than let it stand.

F Clinical operating points and calibration

AUC is threshold-free, but a screening tool is deployed at a threshold. This appendix reports what a reader needs in order to judge that. For each arm the threshold is selected on the *validation* split at a fixed target specificity and then transferred unchanged to the test split, which is the honest simulation of deployment; the achieved test specificity shows whether the threshold survives the shift. Cohort prevalence is 0.4887.

Table 7: Operating points at matched epoch 100. The threshold is chosen on validation to hit the target specificity and applied unchanged to test. ECE is a 15-bin expected calibration error.

policy	target spec.	sens.	spec. (test)	PPV	NPV	Brier	ECE
RANDOM	0.85	0.7701	0.8259	0.8087	0.7899	0.1416	0.0355
RANDOM	0.90	0.7162	0.8794	0.8502	0.7643	0.1416	0.0355
ENVELOPE	0.85	0.7735	0.8351	0.8176	0.7942	0.1364	0.0361
ENVELOPE	0.90	0.7306	0.8807	0.8541	0.7738	0.1364	0.0361
INTENSITY	0.85	0.7851	0.8318	0.8169	0.8020	0.1341	0.0320
INTENSITY	0.90	0.7428	0.8696	0.8448	0.7797	0.1341	0.0320

Two observations. First, the threshold transfers well: a validation target of 0.90 yields test specificity 0.8794 to 0.8807, so the operating point is stable across the split. Second, the AUC advantage survives as a usable sensitivity gain. At a target specificity of 0.90, INTENSITY detects 0.7428 of cases against 0.7162 for the unguided null, a paired difference of +0.0266 (CI [+0.0136, +0.0396], excluding zero). At a target of 0.85 the gain is +0.0150 (CI [+0.0020, +0.0280]). INTENSITY is also the best-calibrated arm (Brier 0.1341 versus 0.1416; ECE 0.0320 versus 0.0355), which matters because a screening score that is used as a probability must be trustworthy as one.

We report this because a +0.011 AUC difference is hard to interpret clinically, whereas “two to three more cases detected per hundred at the same false-positive rate” is not. The same caveats apply as everywhere else in this paper: one dataset, one split, and one pretraining run per policy.

G Broader impact and ethics

This study uses a public, de-identified retinal dataset (FairVision) under its release terms and involves no new data collection and no human subjects. No model reported here is clinically validated or intended for deployment: the frozen-probe AUCs are research measurements on one dataset and one task, and are not evidence of clinical utility.

The subgroup analysis in Section 5.3 is a caution, not a contribution. Every masking policy we tested leaves the same groups worst-served, and the best-performing policy narrows none of them reliably,

even though every subgroup point estimate rises. We would regard deployment of any encoder reported here, without a dedicated and prospective fairness evaluation on the target population, as inappropriate. We also stress that the audit is AUC-only: it contains no subgroup calibration, no sensitivity at a fixed clinical specificity, no predictive-value analysis and no intersectional breakdown, so it should not be mistaken for a complete equity assessment.

Releasing pretrained weights carries the usual dual-use consideration that they may be applied to populations unlike the training cohort. We document the cohort composition (Appendix C) so that such a mismatch is detectable by anyone reusing the weights. The demographic attributes are self-reported categories inherited from the dataset’s coding; they are used only for the subgroup audit and are never model inputs.

H Published ablations on informed versus random masking

Our result should be read against the following published comparisons, which we collected while positioning this work. The pattern is that random masking is a strong baseline and informed masking wins only under some objectives and protocols; the deficit is largest under frozen evaluation.

Table 8: Published ablations comparing a random or simple baseline against an informed or structured masking scheme. Numbers are as reported by the cited papers; metrics differ across rows and are not comparable between rows.

study	baseline	informed variant	interpretation
MAE [He et al., 2022]	random-75: 84.9 FT / 73.5 lin	block-75: 82.8 / 63.9; grid-75: 84.0 / 66.0	random decisively better
SimMIM [Xie et al., 2022]	83.0 top-1	best block 82.7; square 82.6	random modestly better
SemMAE [Li et al., 2022]	random 66.8 lin	within-part 66.5; whole-part 52.9; adaptive 68.7	naive semantics hurt; scheduled mantics help
SemMAE part quality [Li et al., 2022]	no parts 63.7	raw parts 63.6; refined 65.0	an imperfect semantic model a nothing
AutoMAE [Chen et al., 2023a]	MAE 83.26 FT	AutoMAE 83.32	effectively tied
HPM [Wang et al., 2023b]	random 82.49	moderate 82.95; hardest-only 81.40	benefit only with retained randomness
AttMask [Kakogeorgiou et al., 2022]	random 43.4 k -NN	high 43.5; hint 43.6	nearly indistinguishable
AttG-MAE [Sick et al., 2025]	MAE 78.5 @10% labels	guided 78.1	guidance loses in this protocol
SSiT [Huang et al., 2024]	no saliency 79.48 κ	25% mask 77.53; 50% 79.97	dataset-dependent, non-monotonic

Read together, these say that the *direction* of our finding is not new. What we add is a controlled measurement in a setting where the guidance comes from a trained medical segmentation model rather than from attention or difficulty, and where the comparison is run from a shared checkpoint with the masking policy as the only difference.

I Full paired-contrast table

J Full fp32 re-probe

The RANDOM, INTENSITY and ENVELOPE long-horizon probes were originally fitted under the harness default, which enables autocast. To remove any doubt that the headline effect is a precision artifact, we re-ran the complete probe pipeline for those arms with autocast disabled, so that feature extraction, head optimisation and test evaluation are all fp32, under the exact protocol already used by the ANATOMY and COVER arms. The encoder checkpoint is SHA-256 hashed before and after each run and verified unchanged, so no result here can be contaminated by an accidental encoder update.

[fp32 re-probe table pending[†]](#)

K Fine-tuning narrows but does not erase the gap

Under full fine-tuning rather than a frozen probe, INTENSITY reaches 0.8947 (mean-pool head) against RANDOM’s 0.8868, a gap of +0.0079 compared with +0.0109 frozen. Taking each arm’s

Table 9: Primary contrasts: matched epoch *and* matched precision. Δ is a paired difference on identical test cases; the interval is a 10,000-resample paired bootstrap and p is a DeLong test for correlated ROC curves. q is Benjamini–Hochberg over these nine contrasts. Because the test split was inspected repeatedly across this programme, we read all p and q values here as descriptive rather than confirmatory. Every comparison against the null excludes zero.

contrast	epoch	Δ AUC	95% CI	p	q
ENVELOPE — RANDOM	50	+0.0120	[+0.0068, +0.0173]	<0.0001	<0.0001
ENVELOPE — RANDOM	75	+0.0080	[+0.0028, +0.0132]	0.0025	0.0045
ENVELOPE — RANDOM	100	+0.0062	[+0.0010, +0.0113]	0.0199	0.0299
INTENSITY — RANDOM	50	+0.0099	[+0.0051, +0.0149]	<0.0001	0.0001
INTENSITY — RANDOM	75	+0.0113	[+0.0063, +0.0164]	<0.0001	<0.0001
INTENSITY — RANDOM	100	+0.0109	[+0.0057, +0.0160]	<0.0001	<0.0001
INTENSITY — ENVELOPE	50	-0.0020	[-0.0069, +0.0029]	0.4114	0.4114
INTENSITY — ENVELOPE	75	+0.0033	[-0.0013, +0.0079]	0.1491	0.1677
INTENSITY — ENVELOPE	100	+0.0047	[+0.0004, +0.0091]	0.0294	0.0378

best head instead, the gap is +0.0069. The pretraining difference is partially, but not entirely, absorbed by supervised adaptation — consistent with the masking policy shaping the representation rather than merely the optimisation start point. These runs used a different head and schedule from the frozen probes and are never pooled with them.

L Fine-tuned results

Table 10: Full fine-tuning, same test split. Reported separately from frozen probes and never pooled with them.

policy	head	test AUC
INTENSITY	mean-pool	0.894668
INTENSITY	cross-attention	0.893717
INTENSITY	attentive	0.890100
RANDOM	attentive	0.887763
RANDOM	cross-attention	0.887178
RANDOM	mean-pool	0.886756

All masking arms on the SAME 5 slices (volume idx 2652) - one colour per predictor block, production samplers




Figure 4: All six masking policies (columns) on five B-scans (rows) spanning one volume, drawn with the production samplers. Leftmost column is the unmasked B-scan. Each predictor target block has its own colour; the green contour is the anatomy reference. The random, envelope, and cover arms place axis-aligned rectangles; the anatomy arms grow cell-wise targets that follow the retinal band and therefore change shape from slice to slice.




Figure 5: **Left:** race-stratified AUC at epoch 100 with bootstrap intervals. The black subgroup is lowest under every policy, but note that *every* group improves from RANDOM to INTENSITY. **Right:** max-min race gap against aggregate AUC across 20 checkpoints. The fitted slope is positive but does not survive multiplicity correction or branch-level aggregation (Table 4), so we draw no trend conclusion from it.




Figure 6: Single-slice occlusion attribution against native slice position, averaged over all 3000 test volumes, for all three fine-tuned probes. Despite spanning zero to 7.17M probe parameters they converge on the same structure (mean-pool vs cross-attention pool $r=0.94$). The y -axis is *signed* Δlogit : peaks are slices whose removal lowers the logit, the central dip is slices whose removal raises it, and only slices near zero are genuinely unused.




Figure 7: Window occlusion ($W=7$) amplifies the same signal roughly sevenfold and suppresses single-volume noise. For mean-pooled models, single-slice occlusion systematically underestimates attributed signal; the windowed variant recovers about $25\times$ more. We recommend it as the default primitive for this class of model.




Figure 8: The two-peak structure is an OD/OS storage artefact. Clustering per-volume attribution curves into two groups returns near-perfect mirror images of one another, which is what laterality flipping would produce, rather than bilateral anatomy.




Figure 9: Attribution stratified by TP / FN / TN / FP. Errors reproduce the correct pattern at lower amplitude rather than selecting different anatomy.

Occlusion attribution — representative B-scans across the 3 fine-tune probes
Left pane of each row: original B-scan. Right pane: per-patch $\Delta\logit$ overlay.




Figure 10: Per-patch occlusion heat maps on individual B-scans, drawn on a shared scale with the slice mean subtracted so that maps are comparable across volumes. Between 84% and 91% of patches have a 95% bootstrap interval excluding zero on the glaucoma class ($B=500$), so attribution is broadly distributed rather than concentrated in a few patches. Cross-probe agreement is strong at slice level ($r=0.94$) and moderate at patch level (per-volume $r \approx 0.35-0.48$).




Figure 11: The matched-epoch-50 result arranged as a ladder of increasing anatomical specificity. Companion to Figure 3.




Figure 12: Measured mask geometry per policy, 600 held-out slices. The four rectangle policies are closely matched on every axis, so their comparison is near single-variable. The anatomy policy diverges on all four simultaneously — half the masking ratio, $1.6\times$ the context, $0.4\times$ the predictor loss slots — which is why we decline to attribute its deficit to target shape.