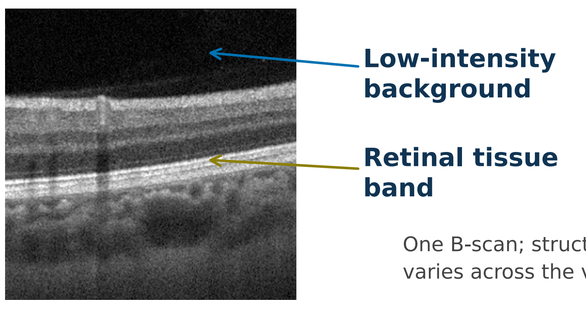
Where to Predict in Retinal OCT? Anatomy-Guided Target Selection for I-JEPA

Yunchu Feng[1] Wai Tak Lau[2]
[1]University of Pennsylvania (yfeng26@engineering.upenn.edu) [2]Columbia University (wl2822@columbia.edu)
GenAI for Health Workshop @ NeurIPS 2026, Sydney

paper, code & pre-registration

OCT structure motivates target selection



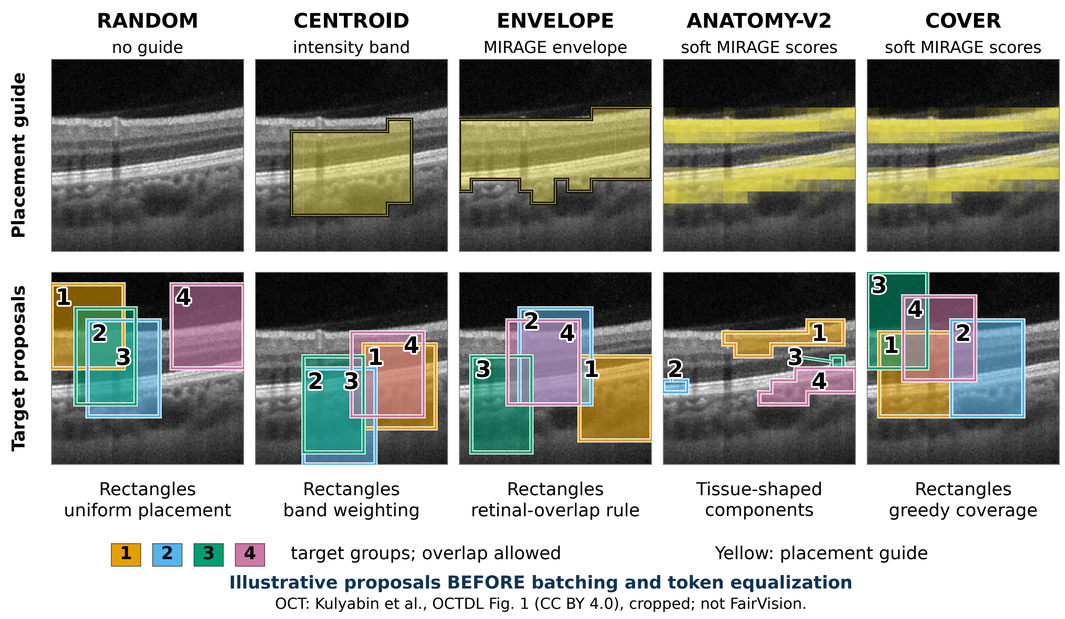
One B-scan; structure varies across the volume.

Question. Can guiding I-JEPA targets toward retinal tissue improve frozen glaucoma features?

Original single runs. Anatomy-guided rectangles scored higher than uniform I-JEPA masking at epoch 50: +0.99 pp with segmentation-free CENTROID, +1.20 pp with ENVELOPE. Target placement and visible context change together.

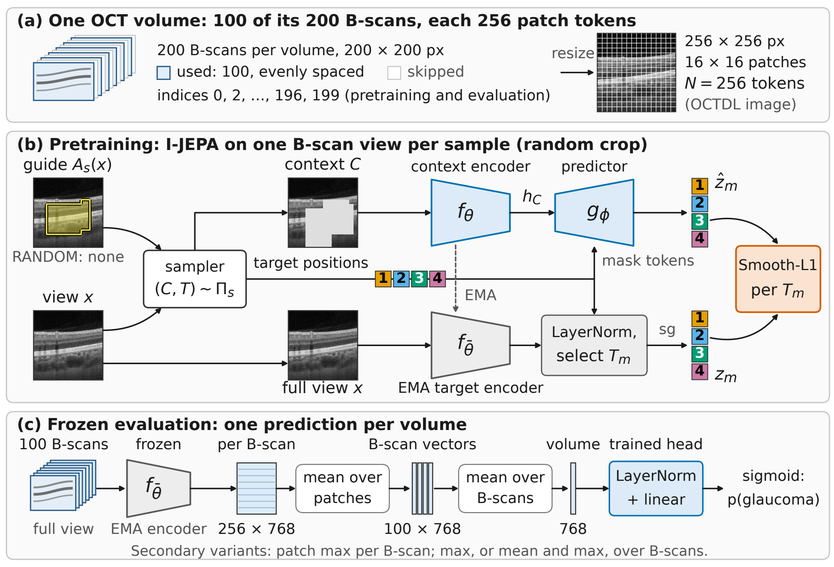
Registered replication: CENTROID *incomplete*, ENVELOPE *incomplete*.

Five target-selection strategies



Proposals before batching; delivered masks can differ (right). Near-uniform pixels can still provide informative latent context.

Same I-JEPA, different sampler; volume-level probe



Only the sampler changes: $(C, T) \sim \Pi_s(\cdot \mid A_s(x))$ with the guide $A_s(x)$; encoder, EMA teacher, predictor and Smooth-L1 loss are shared.

Single runs: anatomy-guided rectangles scored higher

FairVision glaucoma OCT; frozen EMA encoder, mean-pooled features, linear probe; test *N* = 3000 volumes (1466 glaucoma). All continuations start from one epoch-25 uniform I-JEPA checkpoint.

	epoch 50		epoch 75		epoch 100	
strategy	AUC	Δ (pp)	AUC	Δ (pp)	AUC	Δ (pp)
RANDOM	0.8641	—	0.8723	—	0.8746	—
ENVELOPE	0.8761	+1.20	0.8803	+0.80	0.8807	+0.62
CENTROID	0.8740	+0.99	0.8836	+1.13	0.8855	+1.09
ANATOMY-V2	0.8654	+0.13	0.8612	−1.11	—	—
COVER	0.8643	+0.02	0.8639	−0.84	0.8577	−1.68

Δ: minus RANDOM at the same epoch; upper rows fp16 probes, lower rows fp32 (vs an fp32 RANDOM); ANATOMY-V2 stopped at epoch 75. One continuation per strategy: intervals reflect test-case sampling (10,000 paired bootstrap resamples), not retraining.

At epoch 50 (registered endpoint): ENVELOPE +1.20 pp (95% CI [+0.68, +1.73]), CENTROID (no segmenter) +0.99 (95% CI [+0.51, +1.49]); positive at epochs 50, 75 and 100 (ENVELOPE +0.62 to +1.20, CENTROID +0.99 to +1.13).
More anatomy did not add gains: ANATOMY-V2 minus ENVELOPE −1.07 at epoch 50; COVER minus RANDOM +0.02 at epoch 50, −1.68 at epoch 100 (all differences in pp).

Which masks actually reach the model?

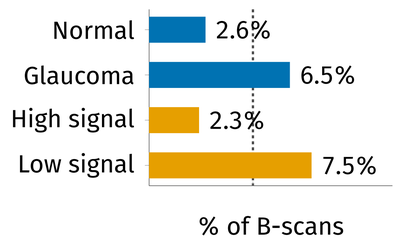
Sampler replay in each trained configuration: 6,400 training views, batch 64, 3 sampler seeds. Tissue is a MIRAGE-derived proxy, not ground truth.

strategy	tissue hidden	target purity	unique targets	loss slots	context (% grid)	AUC ep. 50
RANDOM	53.2%	30.9%	113.1	160.2	27.5 ± 0.1	0.8641
CENTROID	59.7%	38.2%	102.5	160.2	31.5 ± 0.4	0.8740
ENVELOPE	70.7%	42.1%	110.2	160.2	29.2 ± 0.5	0.8761
COVER	73.3%	40.1%	120.1	160.2	25.2 ± 0.4	0.8643
ANATOMY-V2	79.7%	96.9%	54.0	64.0	63.3 ± 0.4	0.8654

- CENTROID and ENVELOPE leave more context visible than RANDOM: +4.0 points (95% CI 3.6–4.4) and +1.6 (1.1–2.1); COVER leaves less.
- ANATOMY-V2 has the purest targets (96.9%) but 64.0 loss slots versus 160.2.

Target location, target size and visible context are coupled.

How reliable are the guides?



CENTROID band vs MIRAGE envelope: gross disagreement in 4.8% of B-scans (dashed; 95% CI 3.7–5.9; all training volumes 5.3%).

Envelope plausibility flags: 3.8% (2.7–5.0).

MIRAGE guide: GOALS foreground Dice 0.925, in-domain only.

Agreement of two automatic guides, not anatomical accuracy (no layer labels). Bars: unweighted rates, stratified sample of 300 volumes (3,000 B-scans); signal: tertiles of an SNR proxy.

Takeaways

- Anatomy-guided rectangles improved frozen glaucoma probes in the original runs, even with a segmentation-free intensity centroid.
- More anatomical specificity (tissue shapes, coverage) added no further gain.
- Target placement, target size and visible context move together. Registered replication labels: CENTROID *incomplete*, ENVELOPE *incomplete*.

Registered replication: repeated continuations and matched-budget control

Repeated continuations Epoch-50 test AUC; originals re-probed, new runs from the epoch-25 checkpoint. Mean and Δ (vs RANDOM, same seed): registered seeds.

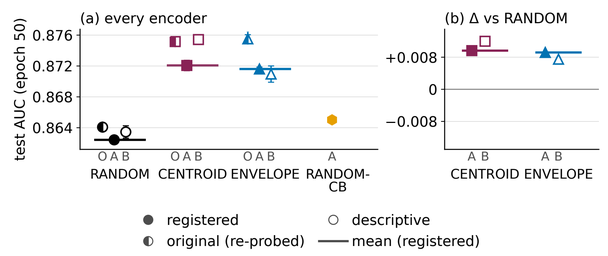
strategy	original	seed A	seed B*	mean	Δ (pp)
RANDOM	0.8641	0.8624	0.8635	0.8624	—
ENVELOPE	0.8755	0.8716	0.8710	0.8716	+0.92
CENTROID	0.8752	0.8721	0.8754	0.8721	+0.96

Matched-budget control RANDOM-CB: uniform placement with CENTROID's per-image budgets; seed A, one run per arm.

contrast	Δ (pp)
CENTROID − RANDOM-CB (placement)	+0.70
RANDOM-CB − RANDOM (budget)	+0.26

Pooling Maxima within or across B-scans; mean Δ (pp) vs RANDOM (descriptive; mean+max: paper).

max	within	across
ENVELOPE	+0.69	+1.84
CENTROID	+1.42	+1.88



Registered labels: CENTROID *incomplete*, ENVELOPE *incomplete*. **Matched-budget control:** not assessed (no registered RANDOM range). *Incomplete at the freeze: descriptive. **Label rule:** consistent direction if every registered Δ > 0 and mean Δ > RANDOM range; reversed if mean Δ ≤ 0; else within run-to-run variation; incomplete with too few on-time runs.

Error bars: s.d. over probe-head seeds. Letters: seed index (O: original). Open markers: descriptive runs (late or incomplete). Panel (b): guided minus RANDOM, same seed, on the original AUC scale. Tables show differences in pp.

Acknowledgments. We thank Prof. Kaveri A. Thakoor for research guidance, medical expertise and retinal anatomy insights, discussions of prior work, and manuscript feedback.

I-JEPA: Assran et al., CVPR 2023. MIRAGE: Morano et al., npj Digit. Med. 2025. FairVision: Luo et al., 2024. All illustrative OCT examples: Kulyabin et al., OCTDL, Fig. 1 (arXiv:2312.08255v3), adapted (cropped, downsampled, overlays added) under CC BY 4.0 (creativecommons.org/licenses/by/4.0); illustration only. No FairVision images are shown.